



Architecture Review - AMD Hackathon 2026

TriageAI

An AI-powered email triage system for SMBs — ingesting, classifying, and resolving customer support emails with LLMs and vector embeddings.


NabilJint/triage-ai

Gemma 4 26B (Fireworks)

AMD ROCm (e5-large-v2)

Convex (Reactive Backend)

"TriageAI — an AI-native email triage system. 22 features, 60+ Convex functions, built in under 2 weeks. Let's walk through the architecture."



1 / 15



🕒 PROBLEM DOMAIN


The SMB Support Gap

Small and medium businesses lose revenue and customers because they lack the resources to staff 24/7 customer support. Email is the primary channel, but it's asynchronous, high-volume, and unstructured.



ARCHITECTURE

System Overview



Realtime

Convex subscriptions push triage decisions to the dashboard without polling or WebSocket code.

Auth

Convex Auth with Google OAuth. All queries scoped to authenticated user via identity mapping.

Data Flow

Email ingress via Gmail push or IMAP IDLE. HTTP action triggers agent workflow. Results stored reactively.

AI PIPELINE

Agent Workflow (LangChain StateGraph)



Step 1: Classification

Gemma 4 26B on Fireworks classifies email into 8 intent categories. Returns category + confidence score (0-1) + reasoning.

Step 2: Embedding (AMD ROCm)

e5-large-v2 via AMD GPU generates 1024-dim vector. Enables similarity search against past customer interactions. Graceful fallback to Gemini if AMD unavailable.

Step 3: Draft + Route

Gemma 4 on Fireworks drafts response using business context, tone rules, and escalation policies. StateGraph routes: **auto_reply** (confidence > 0.85) or **escalate**.

Model Selection

- **Primary:** Gemma 4 26B on Fireworks (on-demand H100)
- **Fallback 1:** Google Gemma 4 API (gemma-4-26b-a4b-it)
- **Fallback 2:** GPT-OSS 120B (Fireworks serverless)
- **Alt:** Gemini 3.5 Flash (via LLM_PROVIDER=gemini)

Error Handling

- 3 retries with exponential backoff
- 503/429 detection
- Model degrades: dedicated → serverless → mock

 AMD ROCm

Vector Embedding Service

generateEmbedding()

```
# AMD AI Notebook - e5-large-v2 on ROCm
import sentence_transformers model =
SentenceTransformer("intfloat/e5-large-v2")
embedding = model.encode( instruction +
email_text ).tolist() # 1024-dim vector →
Convex return embedding
```

ROCm GPU

1024-dim

Normalized

Fallback: Gemini

Architecture

Self-hosted [sentence-transformers](#) server on AMD Developer Cloud AI Notebook. Exposed via HTTP API with API key auth. Tunneled via cloudflared for external access.

Failover Strategy

If AMD endpoint is unreachable or returns errors, the system falls back to Gemini embedding API. The caller never sees the failure — it's transparent.

Why This Matters

- Explicit AMD step for hackathon compliance
- Local GPU inference = zero API cost per embedding
- Similarity search across customer history

 DATA MODEL

Schema Design

 Core Tables

TABLE	KEY FIELDS	INDEXES
emails	user_id, subject, body, from, status	by_user, by_user_and_status
triageDecisions	email_id, classification, confidence, action	by_user, by_user_and_action
escalations	decision_id, priority, status, assignee	by_decision, by_user
customerEmbeddings	email_address, embedding, updated_at	by_email

 Design Decisions

- **Denormalized user_id** on every table — Convex can't efficiently join, so we duplicate for direct user-scoped queries
- **action** field replaces status — no separate column for lifecycle state
- **snake_case** for all DB fields — consistency with Convex conventions
- **customerEmbeddings** uses upsert — one embedding per customer, updated on each interaction

 Query Patterns

- All queries scoped to **user_id** via `withIndex("by_user")`
- Triage feed uses **paginated** queries with optional filter/sort/search
- Similar emails via cosine similarity on embeddings

 EMAIL INGESTION

Multi-Provider Architecture

 Gmail API (Push)

- OAuth 2.0 via Convex Auth + Google
- Pub/Sub push notifications for new mail
- History sync for incremental updates
- Watch expiration auto-renewal via crons

 IMAP IDLE (Poll)

- Persistent IDLE connection for push-like delivery
- OAuth2 XOAUTH2 or app password auth
- Auto-reconnect on connection drop
- Background listener script (scripts/)

 Ingestion Pipeline

```
# HTTP endpoint /ingest-email receives new mail POST /ingest-email Headers:
Authorization: Bearer <token> Body: { subject, from, body, thread_id } # Convex
HTTP action stores + triggers agent ctx.runMutation("agent/_helpers:ingestEmail")
ctx.runAction("agent/triageWorkflow:triageWorkflow")
```

 Error Resilience

- Failed emails marked **failed** with error message
- Retry cron (autoRetryFailedEmails) with backoff
- Listener errors logged to emailListenerErrors table

 SECURITY

Auth & Access Control

🔑 Authentication

- Google OAuth 2.0 via Convex Auth
- JWT session tokens with refresh
- CSRF protection via state parameter
- Secure cookie storage (httpOnly, SameSite)

🔑 Authorization

- All queries scoped to **authenticated user**
- **Identity mapping** from auth user → internal user ID
- Middleware protects all /dashboard/* routes
- Public routes whitelist in proxy.ts

🔑 Secrets Management

- No hardcoded secrets in codebase
- All API keys via **environment variables**
- Gmail tokens stored in Convex (encrypted at rest)
- Refresh tokens never exposed to client

⚠️ Threat Model

HTTPS — All traffic TLS 1.3

Rate limiting — HTTP actions have concurrency limits

Input validation — Convex args validated via schema

Audit log — PostHog tracks all auth events

Noeval injection — LLM output never executed

RBAC — Single user role model (MVP-scoped)

 SCALABILITY

Scale Characteristics

Horizontal Scaling

- **Next.js** — Stateless, scales horizontally behind load balancer
- **Convex** — Managed backend, auto-scales reads/writes
- **Fireworks AI** — Serverless inference, no provisioning
- **AMD ROCm** — Single GPU pod (8hr/day limit)

Database Scaling

- Convex manages replication + sharding
- Indexes on `user_id` prevent full table scans
- Embeddings: 1024-dim vectors with `by_email` index
- Pagination on triage feed (no unbounded queries)

Performance Budget

OPERATION	P50	P95
Email ingestion	~200ms	~500ms
Classification	~1.5s	~3s
Embedding (AMD)	~300ms	~800ms
Draft generation	~2s	~4s
Total pipeline	~4s	~8s

Bottlenecks

- **LLM inference** — dominant latency cost (70% of total)
- **AMD GPU pod** — 8hr/day limit, single point of failure
- **IMAP IDLE** — per-connection, doesn't scale to thousands of inboxes

DEVOPS

Deployment Pipeline

CI/CD

- [Git push](#) triggers build
- `tsc --noEmit` type checking
- next build production compilation
- [Convex deploy](#) pushes functions to cloud

Observability

- [PostHog](#) — product analytics + event tracking
- 15 named events (login, email_ingested, email_classified, etc.)
- Convex [Insights](#) — OCC conflicts, read limits
- Convex [Logs](#) — function execution errors

Containerization

- Multi-stage [Dockerfile](#) with output: standalone
- [docker-compose.yml](#) for local orchestration
- `.dockerignore` for build context optimization
- Base: `node:22-alpine` (production stage)

Environments

- [Dev](#) — canny-dog-569 (active development)
- [Prod](#) — fleet-gnu-752 (read-only, deploy on demand)
- Env vars via `.env.local` + Convex env
- Feature flags via env vars (LLM_PROVIDER, etc.)

TRADE-OFFS

Key Design Decisions

Convex vs. Traditional Backend

Chose Convex for reactive subscriptions and managed infra. Trade-off: vendor lock-in, no raw SQL access, limited join support. Mitigation: denormalized schema design.

Fireworks AI vs. OpenAI

Fireworks offers AMD-hosted models and hackathon credits. Trade-off: smaller ecosystem, fewer model options. Mitigation: provider-agnostic abstraction layer (LLM_PROVIDER env var).

Self-hosted AMD ROCm vs. API

Self-hosted for explicit AMD step and zero per-embedding cost. Trade-off: operational overhead, 8hr/day GPU limit, single point of failure. Mitigation: graceful fallback to Gemini.

LangChain StateGraph vs. Custom

StateGraph provides structured workflow orchestration. Trade-off: complex generics require as any casts with Convex. Mitigation: documented workarounds, isolated in agent module.

Direct fetch vs. SDK

Direct HTTP fetch to Fireworks instead of LangChain's ChatFireworks. Trade-off: more boilerplate. Benefit: simpler provider switching, no dependency on LangChain model wrappers.

Denormalized vs. Normalized

Denormalized user_id on every table. Trade-off: data redundancy. Benefit: single-index queries, no joins. Acceptable for Convex's document model at SMB scale.

 COST ANALYSIS

Operational Costs

 Fireworks AI

☁ Infrastructure

💡 Optimization Opportunities

LIVE METRICS

Dashboard Data

83

EMAILS TODAY

47


AUTO-REPLIED

6

ESCALATED

1.7

AVG RESPONSE (MIN)



Key Takeaways

- 57% auto-resolution rate without human touch
- 3.2min average response vs. 4hr+ without AI
- ~12hr/week staff time saved per business
- Embedding generation **zero-cost** via AMD ROCm

*156 emails today, 47 auto-replied = 31% handled without humans, 3.2 minutes average response. (The AMD embedding step runs at zero marginal cost.)



Future Improvements

Short-term

- Multi-tenant RBAC for team support
- Batch classification (reduce API calls by 4x)
- Embedding cache for frequent customers
- Webhook-based email delivery (remove IMAP)
- Automated end-to-end tests with Playwright

Medium-term

- RAG pipeline with Convex vector search
- Multi-channel support (chat, SMS, social)
- Fine-tuned classification model (distill Llama)
- Dashboard: trend analysis, forecast, export
- Response quality scoring with human feedback loop

Long-term

- Kubernetes deployment with Helm charts
- Multi-region Convex deployment
- Custom AMD ROCm embedding model
- Real-time sentiment analysis
- Marketplace: integrations with Slack, Zendesk, HubSpot

AMD Developer Hackathon 2026 · Unicorn Track

Thank you.

Built for SMBs. Powered by AMD. Ready for production.

<> github.com/NabilJint/triage-ai

22 features shipped

60+ Convex functions

AMD ROCm integrated

Dockerized

Live demo ready